

Contributing to the NLP Search Engine

Ground rules, workflow, and review checklist

Hamza Abdul Karim · F23607046 — May 2026

ABSTRACT

Thank you for considering a contribution. This is a teaching-quality project, so changes that keep the code small and easy to read are especially welcome. The sections below describe the ground rules, how to set up an environment, and the review checklist a pull request should satisfy.

1. Ground Rules

- Small files. Each source file in `src/` stays under 200 lines.
- No frameworks where `stdlib` suffices. Pulling in a heavyweight dependency requires a strong justification.
- Tests stay fast. The `pytest` suite must run in well under one second locally.
- Be kind. See the bundled Code of Conduct.

2. Setting Up

```
git clone https://github.com/Hamza35779/NLP-AI-Search-Engine.git
cd NLP-AI-Search-Engine
python -m venv .venv && source .venv/bin/activate
pip install -e ".[dev]"
pre-commit install
```

3. Workflow

- Open an issue describing the change. Skip this for trivial one-line fixes.
- Create a feature branch: `git checkout -b feat/short-name`.
- Write the code and the matching test(s).
- Run the suite: `pytest -q`.
- Open a pull request against `main` and fill out the template.

4. Pull Request Checklist

- Every modified source file is still under 200 lines.
- `pytest -q` passes locally.
- Public API changes are reflected in the relevant PDF under `docs/`.
- `CHANGELOG.pdf` updated under the Unreleased section if user-visible.
- Commit messages follow Conventional Commits (`feat:`, `fix:`, `docs:...`).